

DistlA: A Cost-Effective Dynamic Impact Analysis for Distributed Programs

Haipeng Cai^{*} and Douglas Thain⁺

**Electrical Engineering and Computer Science, Washington State University*

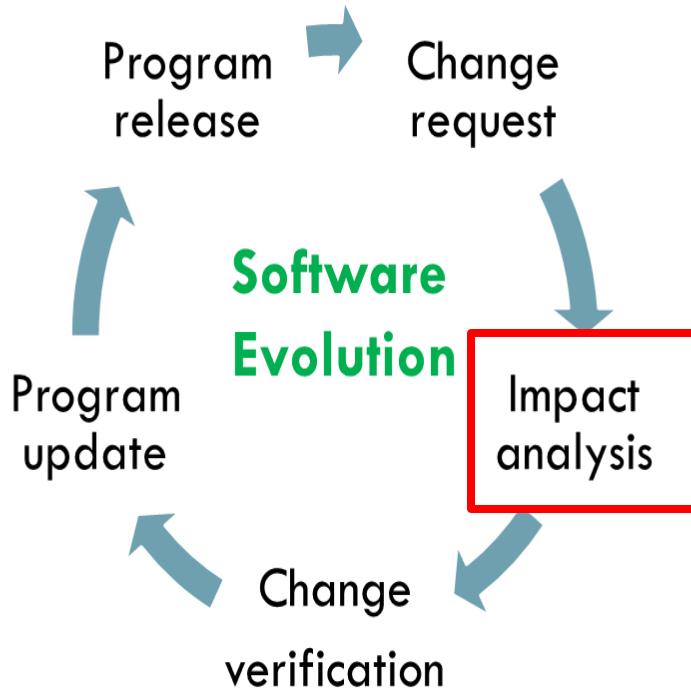
+Computer Science and Engineering, University of Notre Dame



Problem

2

Motivation

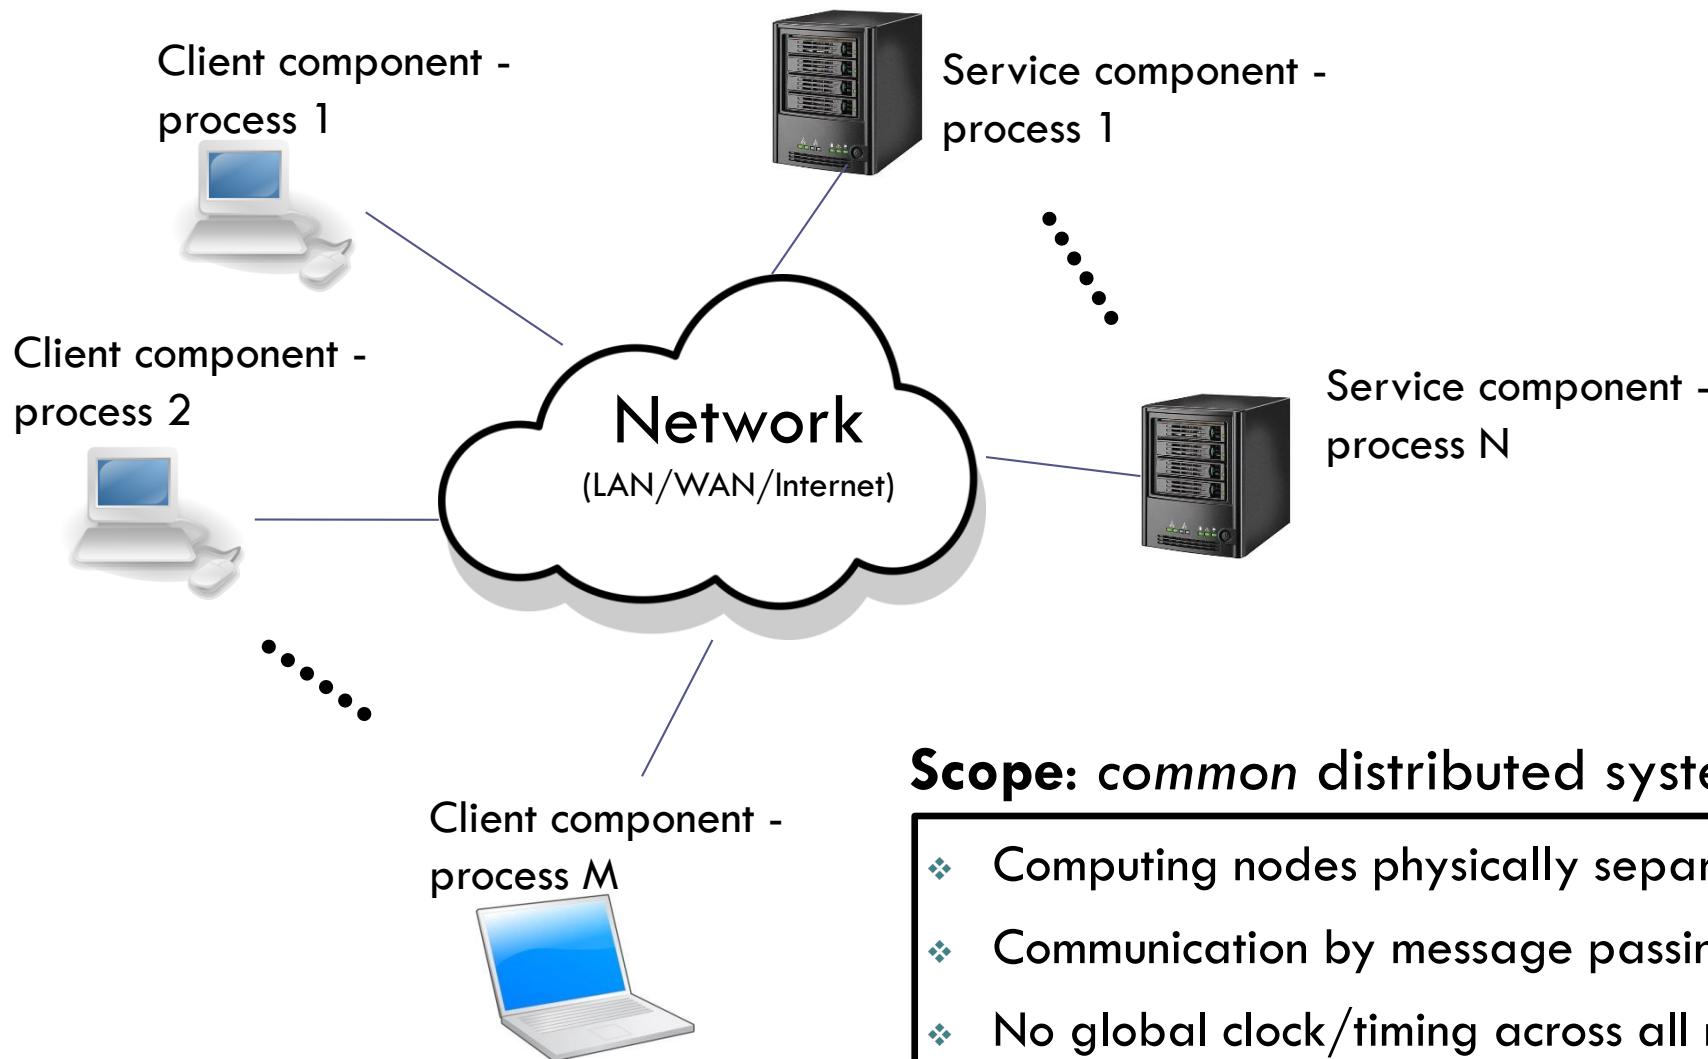


- Various approaches
 - ▣ Mode: *predictive* / *descriptive*
 - ▣ Technique: *static* / *dynamic* / *hybrid* / *repository mining* / *IR* / *coupling*
 - ▣ Granularity: *statement* / *method* / *class* / *file level*
- Different domains
 - ▣ Centralized programs (single/multi-threaded)
 - ▣ Distributed programs (multi-process)
- Predictive, dynamic, method-level
 - ▣ *Proactive*
 - ▣ *More representative of actual behaviors*
 - ▣ *Balanced scalability and precision*
 - ▣ *Not available yet to distributed programs*

Run-time setting of distributed programs

3

Motivation



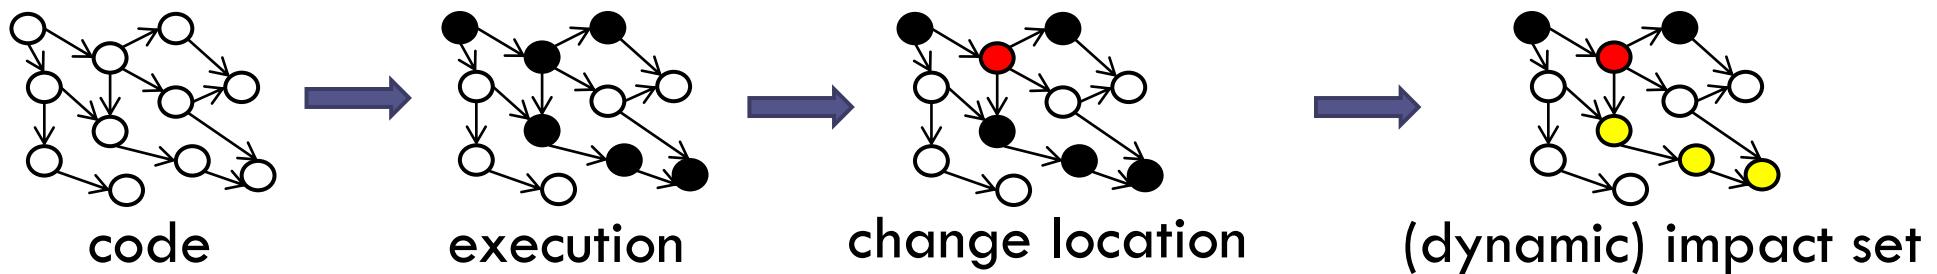
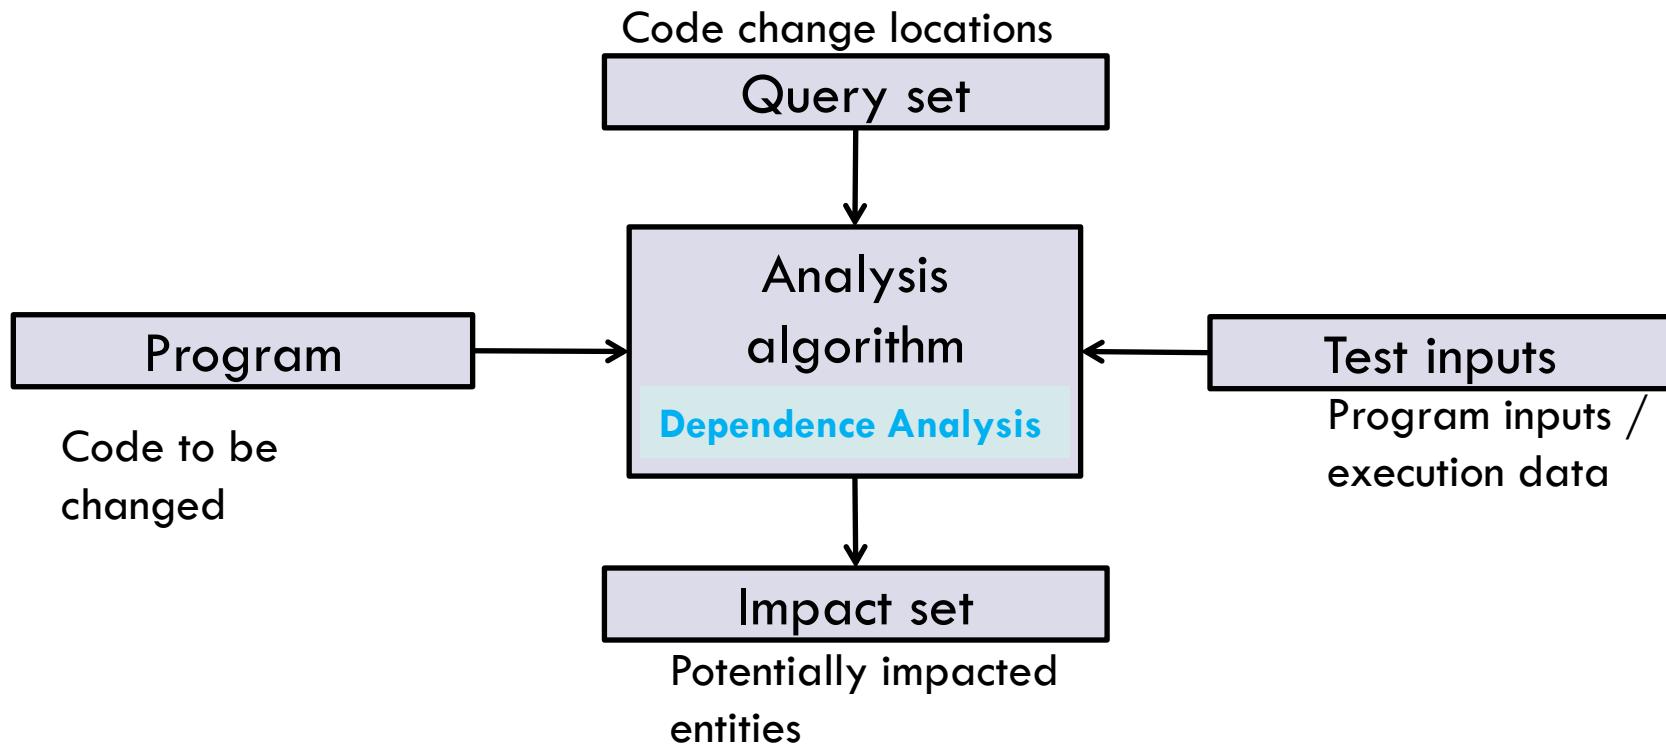
Scope: common distributed systems

- ❖ Computing nodes physically separated
- ❖ Communication by message passing via socket
- ❖ No global clock/timing across all nodes

Dynamic impact analysis

4

Motivation



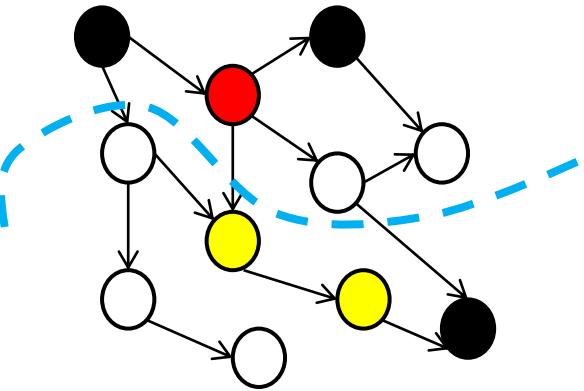
Dependence in distributed programs

5

Motivation

Server

```
1  public class S {  
2      Socket ssock = null;  
3      public S(int port) { ssock = new Socket(port); ssock.accept(); }  
4      char getMax(String s) { ... }  
5      void serve() { String s = ssock.readLine();  
6          char r = getMax(s); ssock.writeChar(r); }  
7      public static int main(String[] a) {  
8          S s = new S(33); s.serve(); return 0; }}
```



Client

```
1  public class C {  
2      Socket csock = null;  
3      public C(String host,int port) { csock = new Socket(host,port); }  
4      void shuffle(String s) { ... }  
5      char compute(String s) { shuffle(s); csock.writeChars(s);  
6          return csock.readChar(); }  
7      public static int main(String[] a) { C c = new C('localhost',33);  
8          System.out.println(c.compute(a[0])); return 0; }}
```

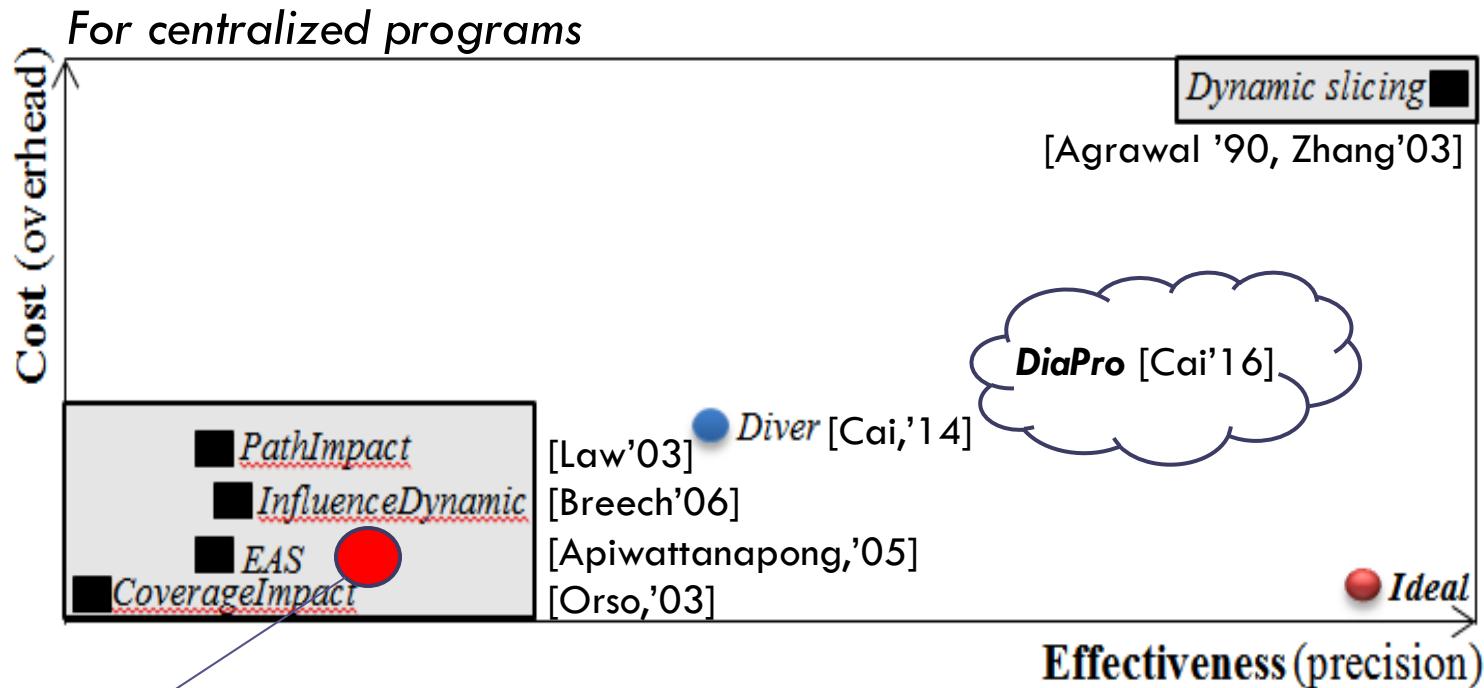
Networking (Socket)

Analysis
algorithm
Dependence Analysis

DistIA: a cost-effective solution

6

Approach



- DistIA
 - Goal
 - cost-effective (rough-yet-rapid [Jackson'00])
 - Strategy
 - *lightweight dynamic dependence approximation*

Dependence approximation

7

Approach

□ Control flow approximation

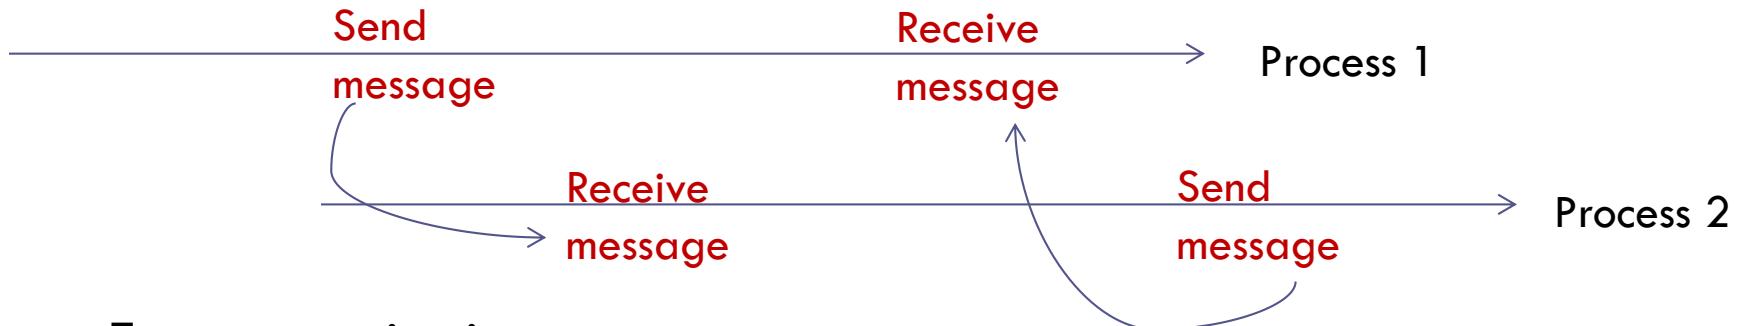
■ Method execution order (partial ordering)



■ Three method events

- Entry, return, and returned-into

- Suffice for single-process partial ordering [Apiwattanapong,'05]



■ Two communication events

- Message sending, and message receiving

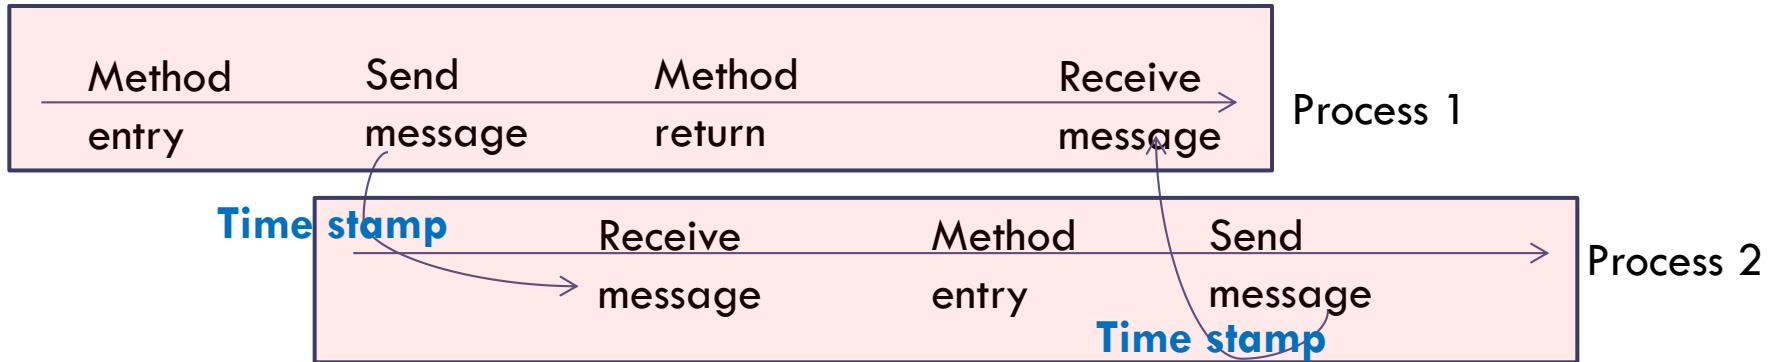
- Necessary for synchronizing the timing of events across processes

Dependence approximation

8

Approach

- Control flow approximation
 - ▣ Global partial ordering [Lamport, '78]



- ▣ Impact inference
 - Happens-before \rightarrow impact relation

$$m1_e \prec m2_x \vee m1_e \prec m2_i \implies m1 \text{ impacts } m2$$

e: entry, x: return, i: returned-into

$$IS(c) = \{m \mid c_e \prec m_i \vee c_e \prec m_x\}$$

IS: impact set, c: query

$$E \prec E' \iff T_L(E') \geq T_F(E')$$

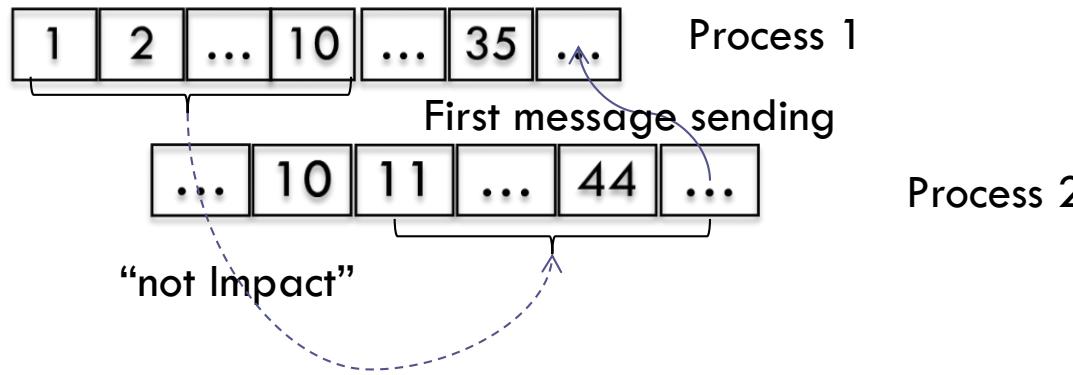
T_F : timestamp of first occurrence, T_L : timestamp of last occurrence

Dependence approximation

9

Approach

- Data flow approximation
 - ▣ Message-passing semantics

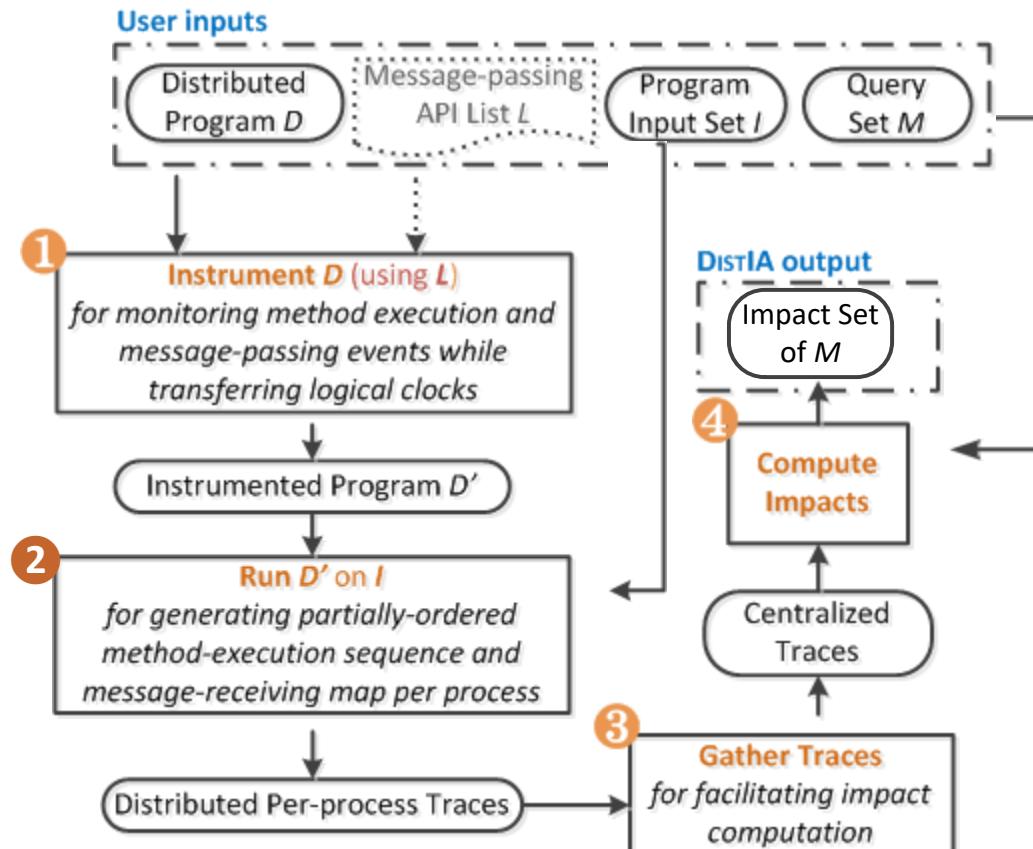


- Control + data flow approximation

$$P_j{}^{m'} \prec P_i{}^m := \begin{cases} T_L(P_i{}^m) \geq T_F(P_j{}^{m'}), & \text{if } i = j \\ T_S(P_j) \neq \text{null} \wedge T_L(P_i{}^m) \geq \max(T_F(P_j{}^{m'}), T_S(P_j)), & \text{if } i \neq j \end{cases}$$

$$IS(c) = \{m \mid c_e \prec m_i \vee c_e \prec m_x\}$$

Workflow



Algorithms

- Communication event monitoring
- Impact computation

Application to real-world distributed programs

11

Evaluation

□ Subject programs

Subject	Description	#SLOC	Test inputs
MultiChat (r5)	C/S chat app, Socket I/O Stream	470	Integration
NIOEcho (r69)	C/S echo service, Java NIO	412	Integration
OpenChord (v1.0.5)	P2P lookup service, hybrid IO	38,084	integration
ZooKeeper (v3.4.6)	Coordination service, hybrid IO	62,450	Integration, system, load
Voldemort (v1.9.6)	Key-value store, hybrid IO	163,601	Integration, system, load
Freenet (v0.7.0)	Anonymous data-sharing, hybrid IO	196,281	integration

□ Implementation

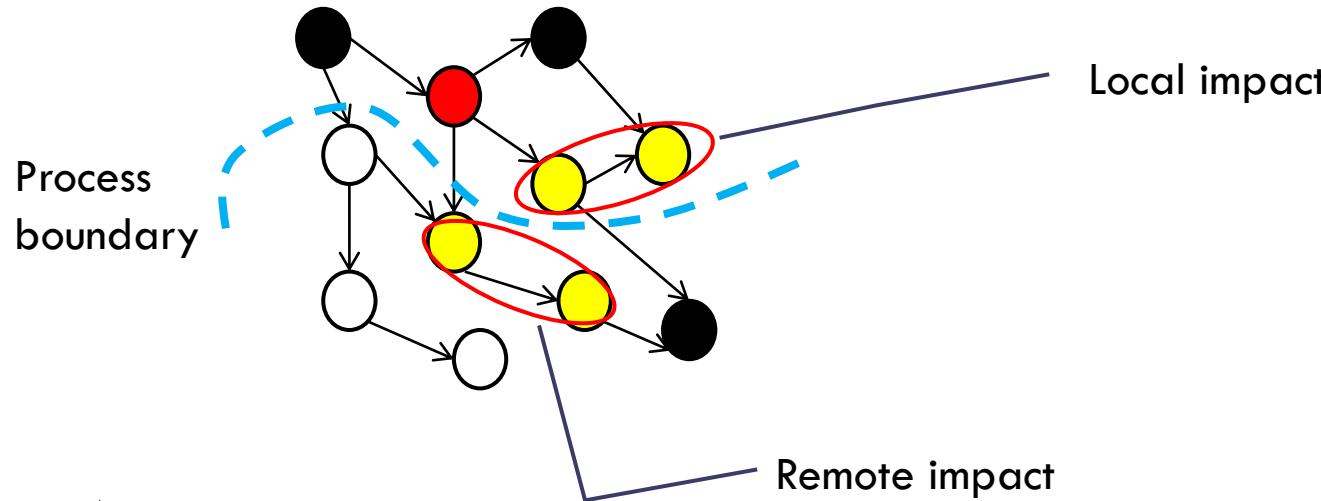
- Non-intrusive instrumentation dealing with a variety of distributed system architectures

Research question: effectiveness

12

Evaluation

- How effective is DistIA compared to “existing options”?
 - ▣ Coverage-based solution (*MCov*) as baseline
- Metrics
 - ▣ Impact set size ratio: DistIA over *Mcov*
 - Assuming both are “dynamically” sound/soundy [Livshits et al., ’15]
 - ▣ Whole impact set (all) and two subsets: local impact set, remote impact set

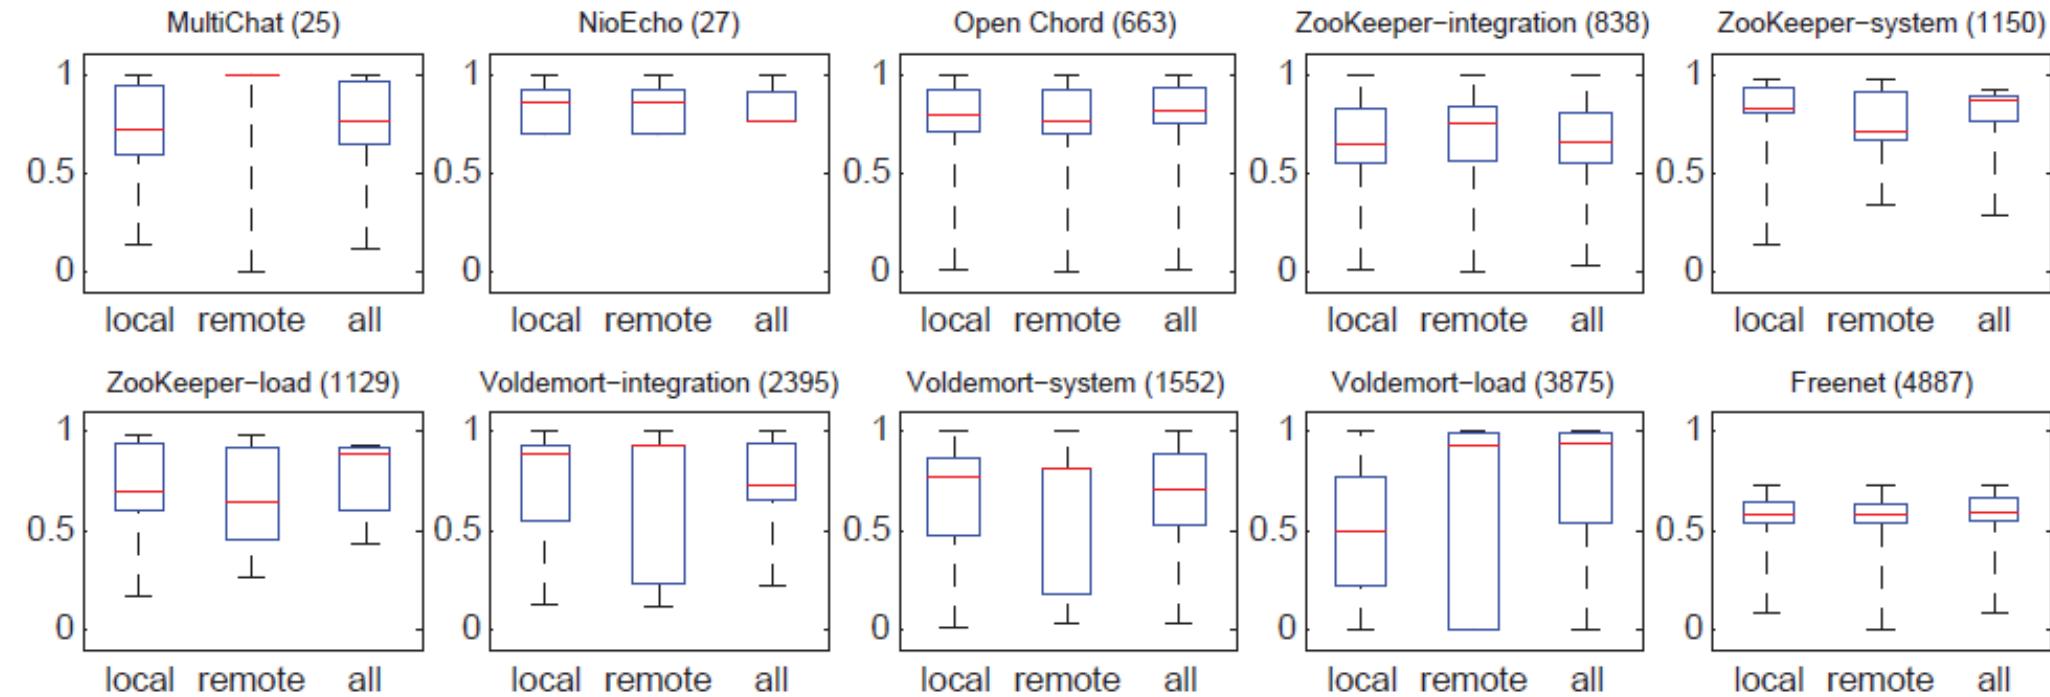


- Two DistIA variants
 - ▣ Basic: control-flow only
 - ▣ Enhanced: data + control flow

Result: effectiveness

13

Evaluation



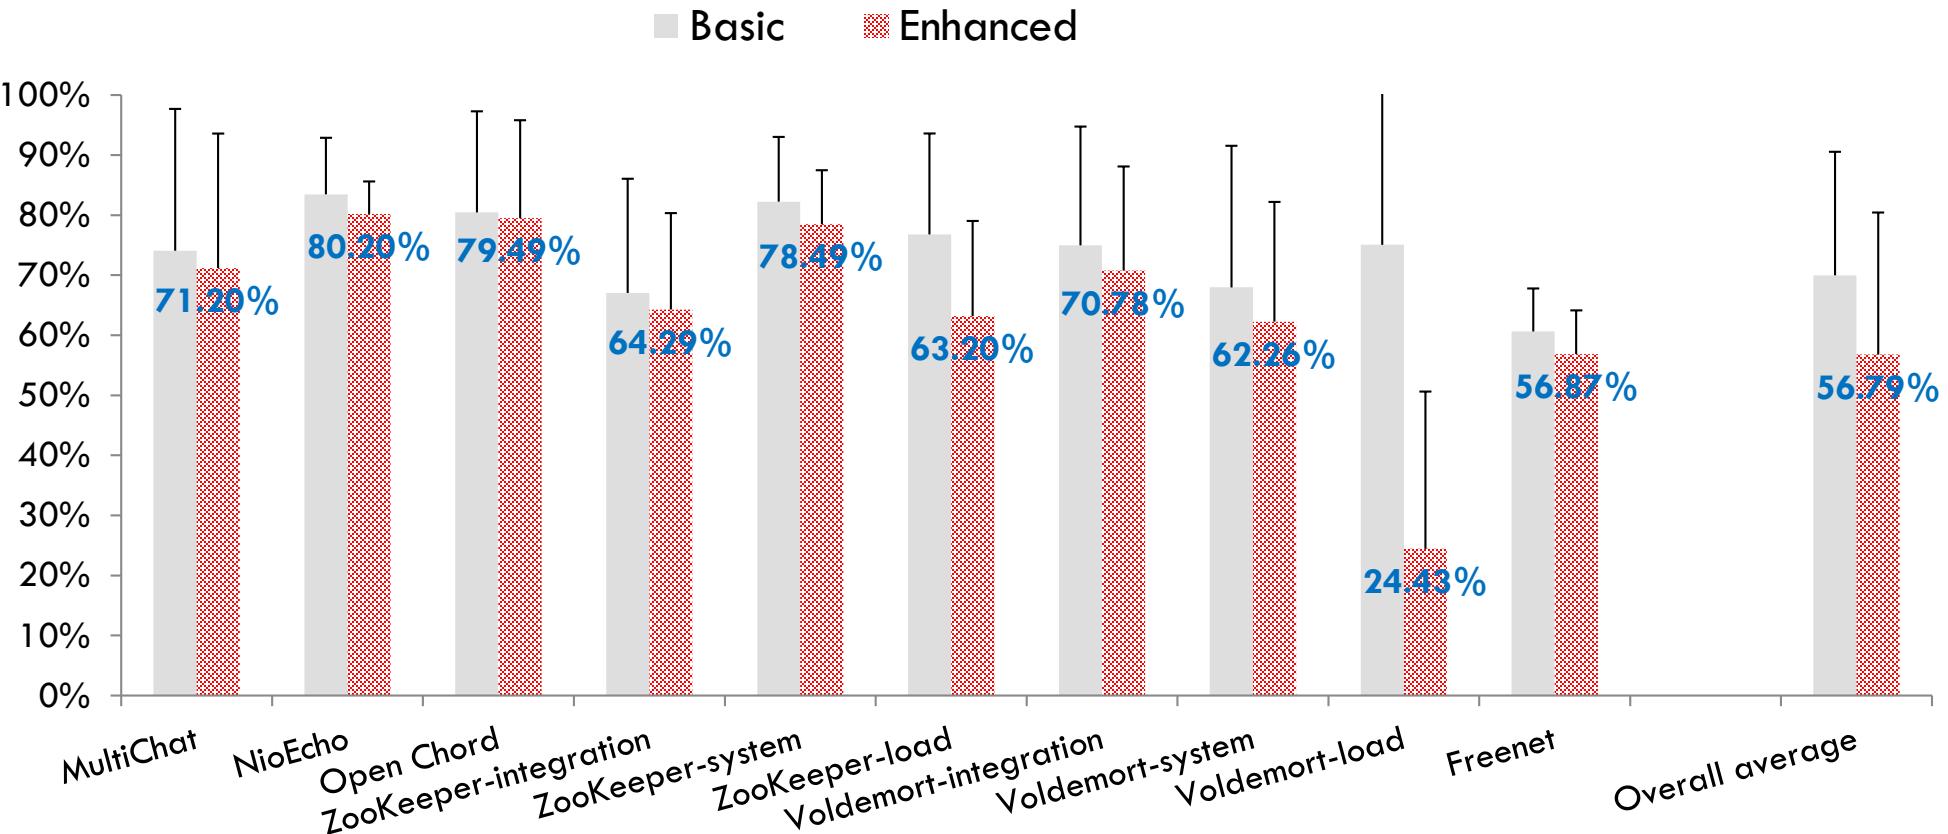
Distribution of impact set size ratios of DistIA-basic/Mcov, the lower the better

Mean impact-set reduction: 31%

Result: effectiveness

14

Evaluation



Mean impact set size ratios of DistIA-basic versus DistIA-enhanced, **the lower the better**

Overall mean impact-set reduction: 43%

Research question: costs

15

Evaluation

- How efficient and scalable would DistIA be?
 - ▣ Practicality of using it in terms of overheads
- Metrics
 - ▣ Time cost
 - ▣ Storage cost
 - ▣ Run-time slowdown

Results: costs

16

Evaluation

Phase	Range	Mean
Instrumentation	12~165 seconds	62 seconds
Run-time slowdown	1~21%	8%
Impact-set querying	4~114 milliseconds	66 milliseconds

Time costs of DistIA enhanced; the basic version costs even less

*Storage cost < 1 MB

Research question: impact distribution

17

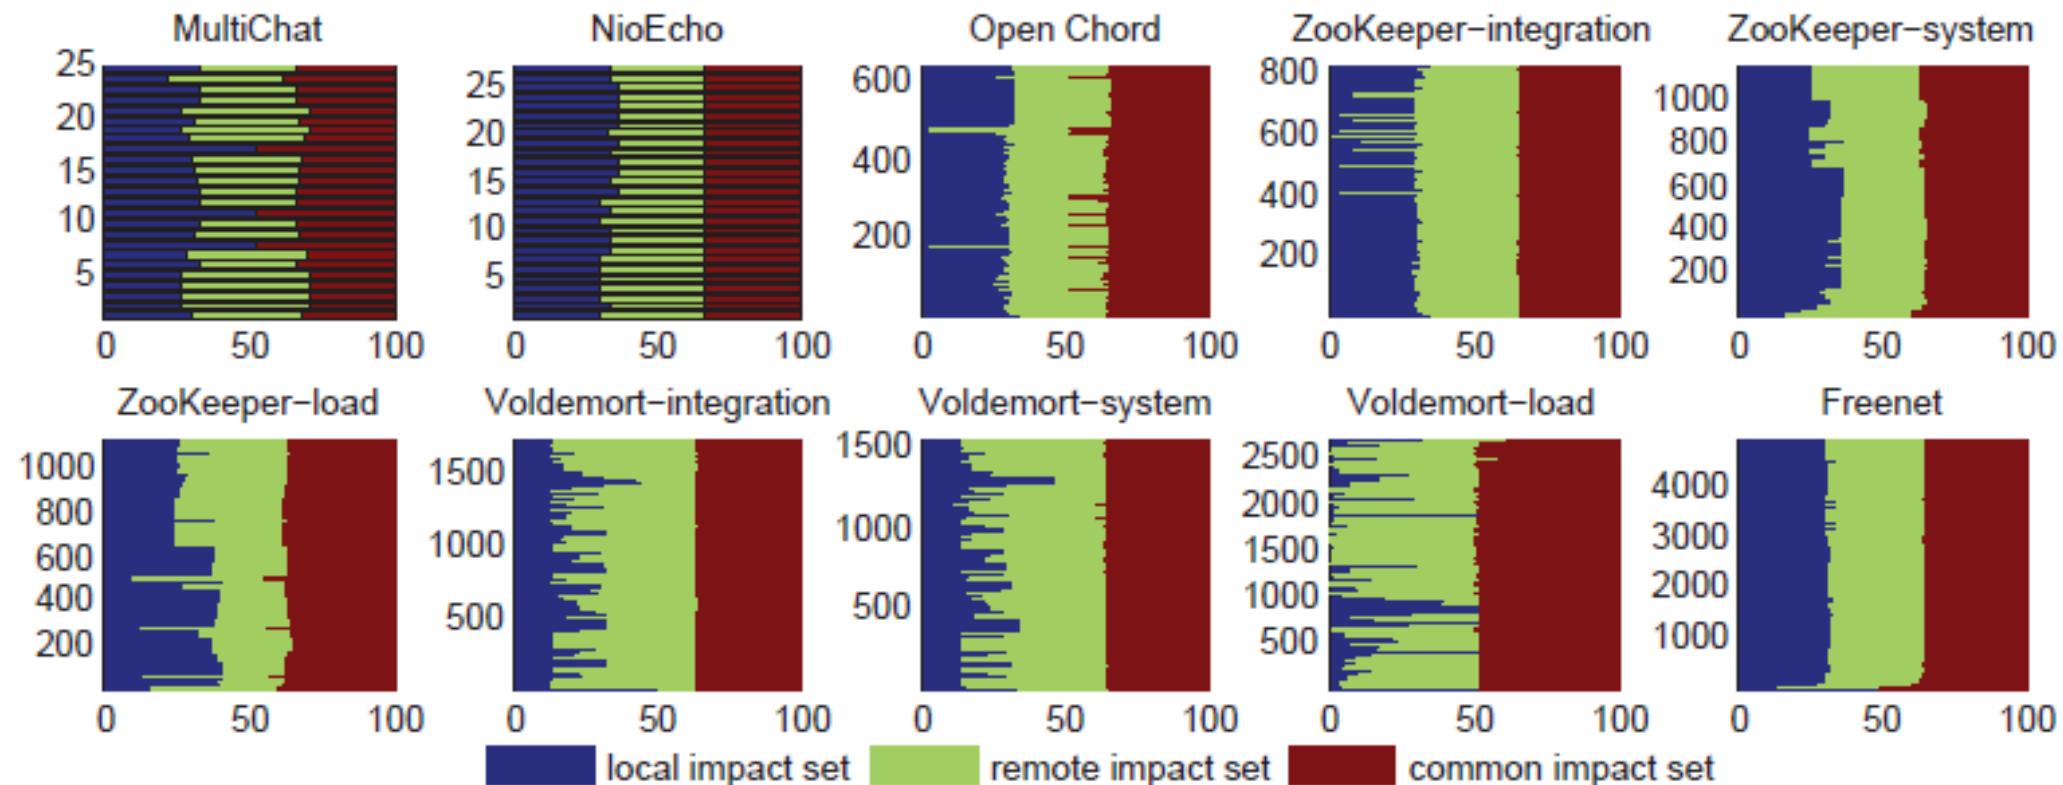
Evaluation

- How are the impacts distributed across process boundaries?
 - ▣ Component-level structure of distributed programs
- Metrics
 - ▣ Impact-set breakdown: local, remote, common impact sets

Result: impact distribution

18

Evaluation



Breakdown of each impact set (given by DistIA-basic) into three **disjoint** subsets

□ Contributions

- The first dynamic impact analysis for common distributed programs with socket-based message passing
- Open-source implementation of the analysis working on real-world, large distributed systems of various architectures
- Empirical evidences showing its promising effectiveness and scalability

□ Future work

- Explore other cost-effectiveness options (with better precision)
- Exploit its use in distributed system testing and understanding

Acknowledgements

20

- ONR grant to Notre Dame and WSU faculty startup fund for financial support
- Anonymous reviewers for very valuable comments
- Your attendance and attention



DistIA: A Cost-Effective Dynamic Impact Analysis for Distributed Programs

Drawing on partial ordering of lightweight dynamic information (method execution events) and simple message-passing semantics heuristics to offer a cost-effective impact-analysis option for real-world distributed programs.

Haipeng Cai

<http://eecs.wsu.edu/~hcai/>

hcai@eecs.wsu.edu

Case study I

22

- How precise are the DistEA impact sets relative to actual dynamic impacts?

Subject & input	DISTEA IS (precision)		Manual true IS		MCov IS	
	local	remote	local	remote	local	remote
MultiChat	1 (100%)	13 (69.2%)	1	9	3	21
MultiChat	13 (76.9%)	2 (50%)	10	1	22	3
Voldemort-system	4 (100%)	23 (56.5%)	4	13	740	809
Voldemort-system	3 (33.3%)	0 (-)	1	0	811	440
Voldemort-load	13 (46.1%)	41 (41.4%)	6	17	288	500
Overall average	6.8 (71.2%)	15.8 (51.7%)	4.7	8	373	354.6

- Overall mean precision
 - DistEA-basic: ~60%
 - DistEA-enhanced: ~70%

Case study 2

23

- How may DistIA results help with understanding inter-process interaction in distributed programs?
 - ▣ NioEcho
 - clearly showing the request initiation from client and server's response by echoing the message received, followed by client's steps in receiving the reply and processing it
 - ▣ ZooKeeper
 - Helped identify the coordination server relays the client request to a worker thread that interacts with database to carry out the client inquiries
 - ▣ In particular: the appearance of communication events in the trace and the timestamp (ordering) of method events are very helpful with sorting out the interactions